


Mutual Exclusion


Nir Shavit
Sub-ing For Nancy Lynch
Distributed Computing
Fall Term

Sequential Computation




2

Concurrent Computation



3

Asynchrony



4

Model Summary

- Multiple *threads*
 - Sometimes called *processes*
- Multiple *CPU's*
 - Sometimes called *processors*
- Single shared *memory*
- *Objects* live in memory
- Unpredictable asynchronous delays

© 2003 Herlihy and Shavit

5


Parallel Primality Testing

- Challenge
 - Print primes from 1 to 10^{10}
- Given
 - Ten-processor multiprocessor
 - One thread per processor
- Goal
 - Get ten-fold speedup (or close)

© 2003 Herlihy and Shavit

6

Load Balancing



- Split the work evenly
- Each thread tests range of 10^9

© 2003 Herlihy and Shavit

7

Procedure for Thread i

```
void run(int i) {
    for (j = i*109+1, j < (i+1)*109; j++) {
        if (isPrime(j))
            print(j);
    }
}
```

© 2003 Herlihy and Shavit

8

Issues


- Larger Num ranges have fewer primes
- Larger numbers harder to test
- Thread workloads
 - Uneven
 - Hard to predict
- Need *dynamic* load balancing

rejected

© 2003 Herlihy and Shavit

9

Shared Counter



© 2003 Herlihy and Shavit

10

Procedure for Thread i

```
int counter = new Counter(1);

void thread(int i) {
    int j = 0;
    while (j < 1010) {
        j = counter.inc();
        if (isPrime(j))
            print(j);
    }
}
```

© 2003 Herlihy and Shavit

11

Procedure for Thread i

```
int counter = new Counter(1);

void thread(int i) {
    int j = 0;
    while (j < 1010) {
        j = counter.inc();
        if (isPrime(j))
            print(j);
    }
}
```

Shared counter object

Stop when every value taken

Increment & return each new value

© 2003 Herlihy and Shavit

12

Counter Implementation

```
public class Counter {  
    private long value;  
  
    public long inc() {  
        return value++;  
    }  
}
```

OK for uniprocessor,
not for multiprocessor

© 2003 Herlihy and Shavit

13

What It Means


```
public class Counter {  
    private long value;  
  
    public long inc() {  
        return value++;  
    }  
}
```

temp = value;
value = value + 1;
return temp;

© 2003 Herlihy and Shavit

14


Uh-Oh



© 2003 Herlihy and Shavit

15

FLP: Facts of Life for Processors



If we could only glue reads and writes...

© 2003 Herlihy and Shavit

16

Challenge

```
public class Counter {  
    private long value;  
  
    public long inc() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

Make these steps
atomic (indivisible)

© 2003 Herlihy and Shavit

17

An Aside: Java™

```
public class Counter {  
    private long value;  
  
    public long inc() {  
        synchronized {  
            temp = value;  
            value = temp + 1;  
        }  
        return temp;  
    }  
}
```

Critical section

Synchronized block

© 2003 Herlihy and Shavit

18

Mutual Exclusion in Detail

- Formal problem definitions
- Solutions for 2 threads
- Solutions for n threads
- Fair solutions
- Inherent costs



© 2003 Herlihy and Shavit

19

Warning

- You will never use these protocols
 - Get over it
- You had better understand them
 - The same issues show up everywhere
 - If you can't reason about these, you won't get far with "real" protocols

© 2003 Herlihy and Shavit

20

Why is Concurrent Programming so Hard?

- Cooking an omelet is easy
- Cooking a five-course meal is hard
- Before we can talk about programs
 - Need a language
 - Describing time and concurrency

© 2003 Herlihy and Shavit

21

Time

- "Absolute, true and mathematical time, of itself and from its own nature, flows equably without relation to anything external." (I. Newton, 1689)
- "Time is Nature's way of making sure that everything doesn't happen all at once." (Anonymous, circa 1970)



© 2003 Herlihy and Shavit

22

Events

- An *event* a_0 of thread A is
 - Instantaneous
 - No simultaneous events




© 2003 Herlihy and Shavit

23

Threads

- A *thread* A is (formally) a sequence a_0, a_1, \dots of events
 - "Trace" model
 - Notation: $a_0 \rightarrow a_1$ indicates order



© 2003 Herlihy and Shavit

24


Example Thread Events

- Assign to shared variable
- Assign to local variable
- Call method
- Return from called method
- Lots of other things ...

© 2003 Herlihy and Shavit

25

Threads are State Machines



© 2003 Herlihy and Shavit

26

States

- Thread State
 - Program counter
 - Local variables
- System state
 - Object fields (shared variables)
 - Union of thread states

© 2003 Herlihy and Shavit

27

Concurrency

- Thread A



- Thread B



© 2003 Herlihy and Shavit

28

Interleavings

- Events of two or more threads
 - Interleaved
 - Not necessarily independent (why?)




© 2003 Herlihy and Shavit

29

Intervals


- An *interval* $A_0 = (a_0, a_1)$ is
 - Time between events a_0 and a_1



© 2003 Herlihy and Shavit

30


Intervals may Overlap



© 2003 Herlihy and Shavit

31

Intervals may be Disjoint




© 2003 Herlihy and Shavit

32

Precedence


Interval A_0 precedes interval B_0



© 2003 Herlihy and Shavit

33

Precedence




- Notation: $A_0 \rightarrow B_0$
- Formally,
 - End event of A_0 before start event of B_0
 - Also called "happens before"

© 2003 Herlihy and Shavit

34

Precedence Ordering




- Remark: $A_0 \rightarrow B_0$ is just like saying
 - 2002 \rightarrow 2003,
 - Middle Ages \rightarrow Renaissance,
- Oh wait,
 - what about this week vs this month?

© 2003 Herlihy and Shavit

35

Precedence Ordering



- Never true that $A \rightarrow A$
- If $A \rightarrow B$ then not true that $B \rightarrow A$
- If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$
- Funny thing: $A \rightarrow B$ & $B \rightarrow A$ might both be false!

© 2003 Herlihy and Shavit

36

Partial Orders

(you may know this already)

- Irreflexive:
 - Never true that $A \rightarrow A$
- Antisymmetric:
 - If $A \rightarrow B$ then not true that $B \rightarrow A$
- Transitive:
 - If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$

© 2003 Herlihy and Shavit

37

Total Orders

(you may know this already)


- Also
 - Irreflexive
 - Antisymmetric
 - Transitive
- Except that for every distinct a, b ,
 - Either $a \rightarrow b$ or $b \rightarrow a$

© 2003 Herlihy and Shavit

38

Repeated Events

```
while (mumble) {  
  a0; a1;  
}
```



© 2003 Herlihy and Shavit

39

Review: Atomic Increment

```
public class Counter {  
  private long value;  
  
  public long inc() {  
    int temp = value;  
    value = value + 1;  
    return temp;  
  }  
}
```

© 2003 Herlihy and Shavit

40

Review: Atomic Increment

```
public class Counter {  
  private long value;  
  
  public long inc() {  
    int temp = value;  
    value = value + 1;  
    return temp;  
  }  
}
```

Allow only one thread at a time

© 2003 Herlihy and Shavit

41

Synchronizaton

```
public interface Lock {  
  public void lock();  
  public void unlock();  
}
```

© 2003 Herlihy and Shavit

42

Synchronizaton

```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```

acquire lock

© 2003 Herlihy and Shavit

43

Synchronizaton

```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```

acquire lock

release lock

© 2003 Herlihy and Shavit

44

Synchronized Atomic Increment

```
public class Counter {  
    private long value;  
    private Lock lock;  
  
    public long getAndIncrement() {  
        lock.lock();  
        int temp = value;  
        value = value + 1;  
        lock.unlock();  
        return temp;  
    }  
}
```

© 2003 Herlihy and Shavit

45

Synchronized Atomic Increment

```
public class Counter {  
    private long value;  
    private Lock lock;  
  
    public long getAndIncrement() {  
        lock.lock();  
        int temp = value;  
        value = value + 1;  
        lock.unlock();  
        return temp;  
    }  
}
```

Acquire Lock

© 2003 Herlihy and Shavit

46

Synchronized Atomic Increment

```
public class Counter {  
    private long value;  
    private Lock lock;  
  
    public long getAndIncrement() {  
        lock.lock();  
        int temp = value;  
        value = value + 1;  
        lock.unlock();  
        return temp;  
    }  
}
```

Acquire Lock

Release Lock

© 2003 Herlihy and Shavit

47

Synchronized Atomic Increment

```
public class Counter {  
    private long value;  
    private Lock lock;  
  
    public long getAndIncrement() {  
        lock.lock();  
        int temp = value;  
        value = value + 1;  
        lock.unlock();  
        return temp;  
    }  
}
```

Critical section

© 2003 Herlihy and Shavit

48

Critical Sections

- Let $CS_i^k \Leftrightarrow$ be thread i 's k -th critical section

© 2003 Herlihy and Shavit

49

Critical Sections


- Let $CS_i^k \Leftrightarrow$ be thread i 's k -th critical section
- And $CS_j^m \Leftrightarrow$ be thread j 's m -th critical section

© 2003 Herlihy and Shavit

50

Critical Sections

- Let $CS_i^k \Leftrightarrow$ be thread i 's k -th critical section
- And $CS_j^m \Leftrightarrow$ be j 's m -th execution
- Then either
 - $\Leftrightarrow \Leftrightarrow$ or $\Leftrightarrow \Leftrightarrow$



© 2003 Herlihy and Shavit

51

Deadlock-Free



- If thread A calls **lock()**
 - And never returns
 - Then other threads must complete **lock()** and **unlock()** calls infinitely often
- System as a whole makes progress
 - Even if individuals starve

© 2003 Herlihy and Shavit

52

Lockout-Free



- If thread A calls **lock()**
 - It will eventually return
- Individual threads make progress
- Exercise:
 - Map deadlock-Free vs lockout-free onto different models of Socialism

© 2003 Herlihy and Shavit

53

Two-Thread vs n -Thread Solutions

- Two-thread solutions first
 - Illustrate most basic ideas
 - Fits on one slide
- Notation watch: for 2-threads
 - Variable i is my thread
 - Variable j is other thread

© 2003 Herlihy and Shavit

54

Two-Thread Conventions

```
public class Thread {  
    private int i;  
    private int j = 1-i;  
  
    public void run() {  
        ...  
    }  
}
```

© 2003 Herlihy and Shavit

55

Two-Thread Conventions

```
public class Thread {  
    private int i;  
    private int j = 1-i;  
  
    public void run() {  
        ...  
    }  
}
```

ID for this thread

© 2003 Herlihy and Shavit

56

Two-Thread Conventions

```
public class Thread {  
    private int i;  
    private int j = 1-i;  
  
    public void run() {  
        ...  
    }  
}
```

ID for this thread

ID for other thread

© 2003 Herlihy and Shavit

57

Two-Thread Conventions

```
public class Thread {  
    private int i;  
    private int j = 1-i;  
  
    public void run() {  
        ...  
    }  
}
```

Henceforth: i is current thread, j is other thread.

© 2003 Herlihy and Shavit

58

Two-Thread Conventions

```
public class Thread {  
    private int i;  
    private int j = 1-i;  
  
    public void run() {  
        ...  
    }  
}
```

Method that does all the work

© 2003 Herlihy and Shavit

59

LockOne

```
public class LockOne implements Lock {  
  
    private bool flag[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

© 2003 Herlihy and Shavit

60

LockOne

```
public class LockOne implements Lock {
    private bool flag[2];
    public void lock() {
        flag[i] = true;
        while (flag[j]) {}
    }
}
```

Set my flag

© 2003 Herlihy and Shavit

61

LockOne

```
public class LockOne implements Lock {
    private bool flag[2];
    public void lock() {
        flag[i] = true;
        while (flag[j]) {}
    }
}
```

Wait for other flag to go false

© 2003 Herlihy and Shavit

62

LockOne Satisfies Mutual Exclusion

- Suppose CS_A concurrent with CS_B
- Before entering critical section
 - $write_A(flag[A]=true) \rightarrow read_A(flag[B]==false)$
 $\rightarrow CS_A$
 - $write_B(flag[B]=true) \rightarrow read_B(flag[A]==false)$
 $\rightarrow CS_B$
- Implications:
 - $read_A(flag[B]==false) \rightarrow write_B(flag[B]=true)$
 - $read_B(flag[A]==false) \rightarrow write_A(flag[B]=true)$

© 2003 Herlihy and Shavit

63

LockOne Satisfies Mutual Exclusion

- Implications:
 - $read_A(flag[B]==false) \rightarrow write_B(flag[B]=true)$
 - $read_B(flag[A]==false) \rightarrow write_A(flag[B]=true)$
- From the code
 - $write_A(flag[A]=true) \rightarrow read_A(flag[B]==false)$
 - $write_B(flag[B]=true) \rightarrow read_B(flag[A]==false)$

© 2003 Herlihy and Shavit

64

LockOne Satisfies Mutual Exclusion

- Implications:
 - $read_A(flag[B]==false) \rightarrow write_B(flag[B]=true)$
 - $read_B(flag[A]==false) \rightarrow write_A(flag[B]=true)$
- From the code
 - $write_A(flag[A]=true) \rightarrow read_A(flag[B]==false)$
 - $write_B(flag[B]=true) \rightarrow read_B(flag[A]==false)$

© 2003 Herlihy and Shavit

65

LockOne Satisfies Mutual Exclusion

- Implications:
 - $read_A(flag[B]==false) \rightarrow write_B(flag[B]=true)$
 - $read_B(flag[A]==false) \rightarrow write_A(flag[B]=true)$
- From the code
 - $write_A(flag[A]=true) \rightarrow read_A(flag[B]==false)$
 - $write_B(flag[B]=true) \rightarrow read_B(flag[A]==false)$

© 2003 Herlihy and Shavit

66

LockOne Satisfies Mutual Exclusion

• Implications:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[B] = \text{true})$

• From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

© 2003 Herlihy and Shavit

67

LockOne Satisfies Mutual Exclusion

• Implications:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[B] = \text{true})$

• From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

© 2003 Herlihy and Shavit

68

LockOne Satisfies Mutual Exclusion

• Implications:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[B] = \text{true})$


• From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$

© 2003 Herlihy and Shavit

69

Cycle!



© 2003 Herlihy and Shavit

70

Deadlock Freedom

• LockOne Fails deadlock-freedom

- Concurrent execution can deadlock

```
flag[i] = true;   flag[j] = true;
while (flag[j]){} while (flag[i]){}

```

- Sequential executions OK

© 2003 Herlihy and Shavit

71

LockTwo

```
public class LockTwo implements Lock {
    private int victim;
    public void lock() {
        victim = i;
        while (victim == i) {};
    }

    public void unlock() {}
}

```

© 2003 Herlihy and Shavit

72

LockTwo

```
public class LockTwo implements Lock {
    private int victim;
    public void lock() {
        victim = i;
        while (victim == i) {}
    }
    public void unlock() {}
}
```

Let other go first

© 2003 Herlihy and Shavit

73

LockTwo

```
public class LockTwo implements Lock {
    private int victim;
    public void lock() {
        victim = i;
        while (victim == i) {}
    }
    public void unlock() {}
}
```

Wait for permission

© 2003 Herlihy and Shavit

74

LockTwo

```
public class Lock2 implements Lock {
    private int victim;
    public void lock() {
        victim = i;
        while (victim == i) {}
    }
    public void unlock() {}
}
```

Nothing to do

© 2003 Herlihy and Shavit

75

LockTwo Claims

- Satisfies mutual exclusion
 - If thread *i* in CS
 - Then **victim == j**
 - Never both 0 and 1
- Not deadlock free
 - Sequential deadlocks
 - Concurrent does not

```
public void lockTwo() {
    victim = i;
    while (victim == i) {}
}
```

© 2003 Herlihy and Shavit

76

Peterson's Algorithm

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim==j) {}
}
public void unlock() {
    flag[i] = false;
}
```

© 2003 Herlihy and Shavit

77

Peterson's Algorithm

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim==j) {}
}
public void unlock() {
    flag[i] = false;
}
```

Announce I'm interested

© 2003 Herlihy and Shavit

78

Peterson's Algorithm

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim==j) {};
}
public void unlock() {
    flag[i] = false;
}
    
```

Annotations for lock():

- Announce I'm interested (points to `flag[i] = true;`)
- Defer to other (points to `while (flag[j] && victim==j) {};`)

© 2003 Herlihy and Shavit

79

Peterson's Algorithm

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim==j) {};
}
public void unlock() {
    flag[i] = false;
}
    
```

Annotations for unlock():

- Wait while other interested & I'm the victim (points to `flag[i] = false;`)

© 2003 Herlihy and Shavit

80

Peterson's Algorithm

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim==j) {};
}
public void unlock() {
    flag[i] = false;
}
    
```

Annotations for lock():

- Announce I'm interested (points to `flag[i] = true;`)
- Defer to other (points to `while (flag[j] && victim==j) {};`)

Annotations for unlock():

- Wait while other interested & I'm the victim (points to `flag[i] = false;`)
- No longer interested (points to the end of the `unlock()` method)

© 2003 Herlihy and Shavit

81

Mutual Exclusion

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == j) {};
}
    
```

- If thread 0 in critical section,
 - flag[0]=true,
 - victim = 1
- If thread 1 in critical section,
 - flag[1]=true,
 - victim = 0

Cannot both be true

© 2003 Herlihy and Shavit

82

Deadlock Free

```

public void lock() {
    ...
    while (flag[j] && victim == j) {};
}
    
```

- Thread blocked
 - only at while loop
 - only if other has the turn
- One or the other must have the turn

© 2003 Herlihy and Shavit

83

Lockout Free

- Thread i blocked only if j repeatedly re-enters so that
 - flag[j] = true and victim == j
- When j re-enters
 - it sets victim to j.
 - So i gets in

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == j) {};
}
public void unlock() {
    flag[i] = false;
}
    
```


© 2003 Herlihy and Shavit

84

The Filter Algorithm for n Threads

There are $n-1$ "waiting rooms" called levels

- At each level
 - At least one enters level
 - At least one blocked if many try
- Only one thread makes it through



© 2003 Herlihy and Shavit

85

Filter

```
class Filter implements Lock {
    int level[n]; // level I want to enter
    int victim[n]; // stop me before I advance again
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i); // busy wait
        }
    }
    public void unlock() {
        level[i] = 0;
    }
}
```

© 2003 Herlihy and Shavit

86

Filter

```
class Filter implements Lock {
    int level[n]; // level I want to enter
    int victim[n]; // stop me before I advance again
    public void acquire(int i) {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i); // busy wait
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

© 2003 Herlihy and Shavit

87

One level at a time

Filter

```
class Filter implements Lock {
    int level[n]; // level I want to enter
    int victim[n]; // stop me before I advance again
    public void acquire(int i) {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i); // busy wait
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

© 2003 Herlihy and Shavit

88

Announce
intention to
enter level L

Filter

```
class Filter implements Lock {
    int level[n]; // level I want to enter
    int victim[n]; // stop me before I advance again
    public void acquire(int i) {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i); // busy wait
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

© 2003 Herlihy and Shavit

89

Give priority to
anyone but me

Filter

Wait as long as someone else is at same or higher level, and I'm designated victim

```
public void acquire(int i) {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while ((∃ k != i) level[k] >= L) &&
            victim[L] == i); // busy wait
    }
}
public void release(int i) {
    level[i] = 0;
}
```

© 2003 Herlihy and Shavit

90

Filter

```
class Filter implements Lock {
    int level[n]; // level I want to enter
    int victim[n]; // stop me before I advance again
    public void acquire(int i) {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists k \neq i$ ) level[k] >= L) &&
                victim[L] == i); // busy wait
        }
    }
}
```


Thread enters level L when it completes the loop

© 2003 Herlihy and Shavit

91

Claim

- Start at level $L=0$
- At most $n-L$ threads enter level L
- Mutual exclusion at level $L=n-1$



© 2003 Herlihy and Shavit

92

Induction Hypothesis

- No more than $n-L+1$ at level $L-1$
- Induction step: by contradiction
- Assume all at level $L-1$ enter level L
- A last to write $victim[L]$
- B is any other thread at level L

```
public void lockO {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while (( $\exists k \neq i$ ) level[k] >= L)
            && victim[L] == i) {}
    }
}
```

© 2003 Herlihy and Shavit

93

First Observation

(1) $write_B(level[B]=L) \rightarrow write_B(victim[L]=B)$

```
public void lockO {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while (( $\exists k \neq i$ ) level[k] >= L)
            && victim[L] == i) {}
    }
}
```

Use the code, Luke!

© 2003 Herlihy and Shavit

94

Second Verse, Same as the First

(2) $write_A(victim[L]=A) \rightarrow read_A(level[B])$

```
public void lockO {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while (( $\exists k \neq i$ ) level[k] >= L)
            && victim[L] == i) {}
    }
}
```

© 2003 Herlihy and Shavit

95

Third Observation



(3) $write_B(victim[L]=B) \rightarrow write_A(victim[L]=A)$

By Hypothesis, A is the last thread to write $victim[L]$

© 2003 Herlihy and Shavit

96

Combining Observations

- (1) $\text{write}_B(\text{level}[B]=L) \rightarrow$ 
- (3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
- (2)  $\rightarrow \text{read}_A(\text{level}[B])$

So A read $\text{level}[B] \neq L$ and could not have entered level L - a contradiction

© 2003 Herlihy and Shavit

97

r -Bounded Waiting

- Want stronger fairness guarantees
- Thread not "overtaken" too much
- Need to adjust definitions

© 2003 Herlihy and Shavit

98

r -Bounded Waiting

- Divide **lock()** method into 2 parts:
 - Doorway interval:
 - Written D_A
 - always finishes in finite steps
 - Waiting interval:
 - Written W_A
 - may take unbounded steps

© 2003 Herlihy and Shavit

99

r -Bounded Waiting

- For threads A and B :
 - If $D_A^k \rightarrow D_B^j$
 - A 's k -th doorway precedes B 's j -th doorway
 - Then $CS_A^k \rightarrow CS_B^{j+r}$
 - A 's k -th critical section precedes B 's $(j+r)$ -th critical section
 - B cannot overtake A by more than r times
- First-come-first-served means $r = 0$.

© 2003 Herlihy and Shavit

100

Fairness Again

- Filter Lock satisfies properties:
 - No one starves (no lockout)
 - But very weak fairness
 - Not r -bounded for any r !
 - That's pretty lame...

© 2003 Herlihy and Shavit

101

Bakery Algorithm

- Basic Idea
 - Take a "number"
 - Wait until lower numbers have been served
- Lexicographic order
 - $(a,b) > (c,d)$
 - If $a > c$, or $a = c$ and $b > d$

© 2003 Herlihy and Shavit

102

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

© 2003 Herlihy and Shavit

103

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];
    Doorway

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

© 2003 Herlihy and Shavit

104

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

Waiting

© 2003 Herlihy and Shavit

105

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];
    I'm interested

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

© 2003 Herlihy and Shavit

106

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];
    Take increasing label

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

© 2003 Herlihy and Shavit

107

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];
    Someone is interested

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i],i) > (label[k],k));
    }
}
```

© 2003 Herlihy and Shavit

108

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i], i) > (label[k], k));
    }
}
```

Someone is interested

With higher label

© 2003 Herlihy and Shavit

109

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    ...

    public void unlock() {
        flag[i] = false;
    }
}
```

© 2003 Herlihy and Shavit

110

Bakery Algorithm

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    ...

    public void unlock() {
        flag[i] = false;
    }
}
```

No longer interested

© 2003 Herlihy and Shavit

111

No Deadlock

- There is always one thread with earliest label
- Ties are impossible (why?)

© 2003 Herlihy and Shavit

112

First-Come-First-Served

- If $D_A \rightarrow D_B$ then A's label is earlier
 - $write_A(label[A]) \rightarrow read_B(label[A]) \rightarrow write_B(label[B]) \rightarrow read_B(flag[A])$
- So B is locked out while flag[A] is true

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i], i) > (label[k], k));
    }
}
```

© 2003 Herlihy and Shavit

113

Mutual Exclusion

- Suppose A and B in CS together
- Suppose A has earlier label
- When B entered, it must have seen
 - flag[A] is false, or
 - label[A] > label[B]

```
class Bakery implements Lock {
    boolean flag[n];
    int label[n];

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i], i) > (label[k], k));
    }
}
```

© 2003 Herlihy and Shavit

114

Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$
- Which contradicts the assumption that A has an earlier label

© 2003 Herlihy and Shavit

115

Bakery Y2³²K Bug

```
class Lock5 implements Lock {
    boolean flag[n];
    int label[n];

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i], i) > (label[k], k));
    }
}
```

© 2003 Herlihy and Shavit

116

Bakery Y2³²K Bug

```
class Lock5 implements Lock {
    boolean flag[n];
    int label[n];
    // FCFS breaks if label[i] overflows

    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n])+1;
        while (∃k flag[k]
            && (label[i], i) > (label[k], k));
    }
}
```

© 2003 Herlihy and Shavit

117

Does Overflow Actually Matter?

- Yes
 - Y2K
 - 18 January 2038 (Unix `time_t` rollover)
 - 16-bit counters
- No
 - 64-bit counters
- Maybe
 - 32-bit counters

© 2003 Herlihy and Shavit

118

Timestamps

- Label variable is really a timestamp
- Need ability to
 - Read others' timestamps
 - Compare them
 - Generate a later timestamp
- Can we do this without overflow?

© 2003 Herlihy and Shavit

119

The **Bad** News

- One can construct a
 - Wait-free (no mutual exclusion)
 - Concurrent
 This part is hard
 - Timestamping system
 - That never overflows

© 2003 Herlihy and Shavit

120


Instead ...

- We construct a Sequential timestamping system
 - Same basic idea
 - But simpler
- Uses mutex to read & write atomically
- No good for building locks
 - But useful anyway

© 2003 Herlihy and Shavit

121

Precedence Graphs




- Timestamps form directed graph
- Edge x to y
 - Means x is later timestamp
 - We say x dominates y

© 2003 Herlihy and Shavit

122

Unbounded Counter Precedence Graph




- Timestamping = move tokens on graph
- Atomically
 - read others' tokens
 - move mine
- Ignore tie-breaking for now

© 2003 Herlihy and Shavit

123


Unbounded Counter Precedence Graph



© 2003 Herlihy and Shavit

124


Unbounded Counter Precedence Graph



© 2003 Herlihy and Shavit

125


Two-Thread Bounded Precedence Graph



© 2003 Herlihy and Shavit

126


Two-Thread Bounded
Precedence Graph



© 2003 Herlihy and Shavit

127

Two-Thread Bounded
Precedence Graph T^2




and so on ...

© 2003 Herlihy and Shavit

128


Three-Thread Bounded
Precedence Graph?



© 2003 Herlihy and Shavit

129


Graph Composition



© 2003 Herlihy and Shavit

130


Three-Thread Bounded
Precedence Graph T^3



© 2003 Herlihy and Shavit

131

Three-Thread Bounded
Precedence Graph T^3




© 2003 Herlihy and Shavit

132

In General

$$T^k = T^2 * T^{k-1}$$

K threads need 3^k nodes



© 2003 Herlihy and Shavit

133

Deep Philosophical Question

- The Bakery Algorithm is
 - Succinct,
 - Elegant, and
 - Fair.
- Q: So why isn't it practical?
- A: Well, you have to read N distinct object fields

© 2003 Herlihy and Shavit

134

Theorem


At least N multi-reader/single-writer registers are needed to solve deadlock-free mutual exclusion.

© 2003 Herlihy and Shavit

135

Proof

Each thread must write to some register



Can't tell whether A is in critical section

© 2003 Herlihy and Shavit

136

Upper Bound

- You need at least N MRSW registers
- Bakery algorithm
 - Uses $2N$ MRSW registers
- So the bound is (pretty) tight
- But what if we use MRMW registers?
 - Like the Filter algorithm?

© 2003 Herlihy and Shavit

137

Bad News Theorem

At least N multi-reader/multi-writer registers are needed to solve deadlock-free mutual exclusion.

© 2003 Herlihy and Shavit

138


Let Prove: Theorem

Deadlock-free mutual exclusion for 3 threads requires at least 3 multi-reader multi-writer fields

© 2003 Herlihy and Shavit

139

Covering State

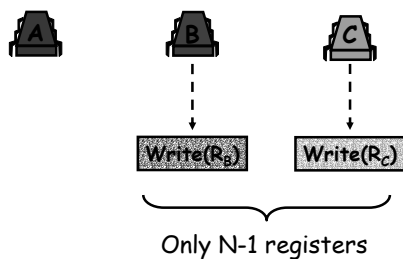


- All registers about to be written
- CS looks empty to all threads

© 2003 Herlihy and Shavit

140


Proof: Assume



© 2003 Herlihy and Shavit

141


Solo Execution



© 2003 Herlihy and Shavit

142

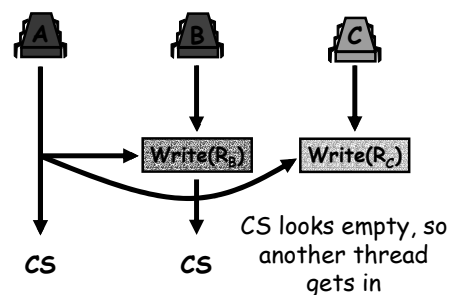
Covering State



© 2003 Herlihy and Shavit

143

Mutual Exclusion Fails



© 2003 Herlihy and Shavit

144

Proof Strategy

- Proved: In a covering state, you need 3 distinct fields
- Claim: a covering state is reachable from any state where CS is empty

© 2003 Herlihy and Shavit

145

Covering State for One Register




Write(R_B)

B has to write to some register to enter CS, so stop it just before

© 2003 Herlihy and Shavit

146

Covering State




- If we run B through CS 3 times, B must return twice to some register, say R_B

© 2003 Herlihy and Shavit

147

Covering State




- Start with B covering register R_B
- Run A until it is about to write to uncovered R_A
- Are we done?

© 2003 Herlihy and Shavit

148

Covering State




- A could have written to R_B
- CS no longer looks empty to some thread

© 2003 Herlihy and Shavit

149

Covering State




- Run B obliterating traces of A in register R_B
- Run B again until it is about to write to R_B
- Now we are done

© 2003 Herlihy and Shavit

150

Inductively We Can Show



- There is a covering state
 - Where k threads not in CS
 - Cover k distinct registers
 - $k=N-1$ delivers proof

© 2003 Herlihy and Shavit

151

Mutual Exclusion in Practice

- Shared FIFO queue
- Written in standard Java™

© 2003 Herlihy and Shavit

152

Mutual Exclusion in Practice

- Shared FIFO queue
- Written in standard Java™

© 2003 Herlihy and Shavit

153

Lock-Based Queue

```
public class Queue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public synchronized void enq(Item x) {
        while (this.tail - this.head == QSIZE)
            this.wait();
        this.items[this.tail++ % QSIZE] = x;
        this.notifyAll();
    }
    ...
}
```

© 2003 Herlihy and Shavit

154

Lock-Based Queue

```
public class Queue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public synchronized void enq(Item x) {
        while (this.tail - this.head == QSIZE)
            this.wait();
        this.items[this.tail++ % QSIZE] = x;
        this.notifyAll();
    }
    ...
}
```

Acquire lock on entry,
release on exit

© 2003 Herlihy and Shavit

155

Lock-Based Queue

```
public class Queue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public synchronized void enq(Item x) {
        while (this.tail - this.head == QSIZE)
            this.wait();
        this.items[this.tail++ % QSIZE] = x;
        this.notifyAll();
    }
    ...
}
```

If Queue is full, release lock,
sleep, try again

© 2003 Herlihy and Shavit

156

Lock-Based Queue

```
public class queue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public synchronized void enq(Item x) {
        while (this.tail - this.head == QSIZE)
            this.wait();
        this.items[this.tail++ % QSIZE] = x;
        this.notifyAll();
    }
}
```

Append the item to the queue

© 2003 Herlihy and Shavit

157

Lock-Based Queue

```
public class Queue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public synchronized void enq(Item x) {
        while (this.tail - this.head == QSIZE)
            this.wait();
        this.items[this.tail++ % QSIZE] = x;
        this.notifyAll();
    }
}
```

Wake up sleeping dequeuers

© 2003 Herlihy and Shavit

158

Observations

- Each method locks entire queue
- No concurrency between methods
- Is this really necessary?

No

And thereby hangs a tale ...

© 2003 Herlihy and Shavit

159

Lock-Free Queue

- Imagine two threads
 - One enqueues only
 - One dequeues only
- Do they need mutual exclusion?

© 2003 Herlihy and Shavit

160

Lock-Free Queue

```
public class LockFreeQueue {
    int head = 0, tail = 0;
    Object[QSIZE] items;

    public void enq(Item x) {
        while (tail - head == QSIZE) {}
        items[tail % QSIZE] = x; tail++;
    }

    public Item deq() {
        while (tail == head) {}
        Item item = items[head % QSIZE]; head++;
        return item;
    }
}
```

© 2003 Herlihy and Shavit

161

Lock-Free Queue

```
public class LockFreeQueue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public void enq(Item x) {
        while (tail - head == QSIZE) {}
        items[tail % QSIZE] = x; tail++;
    }
    public Item deq() {
        while (tail == head) {}
        Item item = items[head % QSIZE];
        head++; return item;
    }
}
```

© 2003 Herlihy and Shavit

162

Lock-Free Queue

```
public class LockFreeQueue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public void enq(Item x) {
        while (tail-head == QSIZE) {};
        items[tail % QSIZE] = x; tail++;
    }
    public Item deq() {
        while (tail == head) {} // Spin while queue is full
        Item item = items[head % QSIZE];
        head++; return item;
    }
}
```

© 2003 Herlihy and Shavit

163

Lock-Free Queue

```
public class LockFreeQueue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public void enq(Item x) {
        while (tail-head == QSIZE) {};
        items[tail % QSIZE] = x; tail++;
    }
    public Item deq() {
        // Put object in queue
        while (tail == head) {}
        Item item = items[head % QSIZE];
        head++; return item;
    }
}
```

© 2003 Herlihy and Shavit

164

Lock-Free Queue

```
public class LockFreeQueue {
    int head = 0, tail = 0;
    Item[QSIZE] items;
    public void enq(Item x) {
        while (tail-head == QSIZE) {};
        items[tail % QSIZE] = x; tail++;
    }
    public Item deq() {
        // Increment tail
        while (tail == head) {} // counter
        Item item = items[head % QSIZE];
        head++; return item;
    }
}
```

© 2003 Herlihy and Shavit

165

Vive La Différence

- The lock-based Queue
 - Is coarse-grained synchronization
 - Critical section is entire method
- The lock-free Queue
 - Is fine-grained synchronization
 - Critical section is single machine instruction

© 2003 Herlihy and Shavit

166

Critical Sections

- Easy way to implement concurrent objects
 - Take sequential object
 - Make each method a critical section
- Like synchronized methods in Java™
- Problems
 - Blocking
 - No concurrency

© 2003 Herlihy and Shavit

167

Amdahl's Law

Sequential fraction

Speedup = $\frac{1}{1 - c + \frac{c}{n}}$

Parallel fraction

Number of processors

Diagram description: The equation $\text{Speedup} = \frac{1}{1 - c + \frac{c}{n}}$ is shown. A box labeled '1' is above the denominator. A box labeled 'c' is above the fraction $\frac{c}{n}$. A box labeled 'n' is below the fraction $\frac{c}{n}$. Arrows point from the text labels to the corresponding parts of the equation: 'Sequential fraction' to the numerator '1', 'Parallel fraction' to the 'c' in the denominator, and 'Number of processors' to the 'n' in the denominator.

© 2003 Herlihy and Shavit

168

Example

- Ten processors
- 60% concurrent, 40% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=2.17= \frac{1}{1 - 0.6 + \frac{0.6}{10}}$$

© 2003 Herlihy and Shavit

169

Example

- Ten processors
- 80% concurrent, 20% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=3.57= \frac{1}{1 - 0.8 + \frac{0.8}{10}}$$

© 2003 Herlihy and Shavit

170

Example

- Ten processors
- 90% concurrent, 10% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=5.26= \frac{1}{1 - 0.9 + \frac{0.9}{10}}$$

© 2003 Herlihy and Shavit

171

Example

- Ten processors
- 99% concurrent, 01% sequential
- How close to 10-fold speedup?

$$\text{Speedup}=9.17= \frac{1}{1 - 0.99 + \frac{0.99}{10}}$$

© 2003 Herlihy and Shavit

172

The Moral

- Granularity matters
 - Long critical sections vs atomic machine instructions
 - Smaller the granularity, greater the speedup

© 2003 Herlihy and Shavit

173

Mutual Exclusion

Nir Shavit
Multiprocessor Synchronization
Fall 2003